

Functional Trust Regions (FTR): A Lagrangian Framework for Stability-Constrained Continual Learning

Kavya Bhand Aadi Joshi

Department of Computer Engineering, Vishwakarma Institute of Technology, Pune 411037, India
{kavya.bhand0806, toaadijoshi}@gmail.com

Abstract

A prior study of Functional Trust Regions (FTR), a method that enforces explicit KL-divergence constraints on functional drift during sequential task learning, reported that the critical stability budget ε^* at which forgetting turns catastrophic is constant across eight architectures, concluding it reflects task structure rather than model geometry. We show this is incomplete: because the dual-ascent optimizer resets its Lagrange multiplier at every task boundary, ε^* is mechanically co-determined by the optimizer’s own hyperparameters, not only by task or architecture. Holding architecture and task fixed, the dual learning rate, the initial multiplier, and the steps-per-task budget each move ε^* by more than an order of magnitude; across six single-factor perturbations, ε^* spans roughly 1 to 90 on the *same* model and task, exceeding the effect of the original 24-fold architecture range. This three-dimensional sensitivity is not irreducible: it collapses substantially onto a single ratio, $S = \lambda_{\text{init}}/(\eta_{\lambda}N)$. Holding S fixed while independently varying its three components by up to $5\times$ keeps ε^* within a factor of 2–3, whereas changing S itself by the same factor collapses ε^* by $4\text{--}11\times$, turning an intractable three-hyperparameter sensitivity into a single number a practitioner can report and match. Holding the schedule fixed and expanding the architecture zoo to 30 models across five representational families (CNN, ResNet, MLP, Vision Transformer, MLP-Mixer), a hierarchical partial-pooling model, replacing the original F-test whose own code and reported degrees of freedom disagreed, initially estimates real, non-negligible between-architecture structure. Applying our own diagnostic standard to our own zoo design, however, we trace most of that structure to two further, independently-diagnosed data-quality issues: a self-inflicted instance of the same schedule confound affecting four architectures, and a curve-fitting artifact affecting two more. Correcting both cuts the model’s between-architecture variance by more than half (from $\tau = 2.13$ to 1.03) and gives a corrected population-level crossover of $\varepsilon^* = 11.07 \pm 0.25$. At this larger, adequately powered sample, the four purely curvature-based statistics we test (Hessian trace, Fisher trace, spectral norm, gradient norm) remain uncorrelated with ε^* , confirming the original curvature-irrelevance claim on grounds its own underpowered $n = 8$ test could not support, though model size shows a small trend worth further testing. Cross-method comparison on a matched architecture subset shows LwF exhibits the same qualitative universal crossover while EWC and SI show none, and further robustness checks (KL direction, task ordering, class-incremental evaluation, task granularity) both support and further qualify the picture. A pretrained ViT-B/16 with a LoRA adapter on Split CIFAR-100, a setting closer to how continual learning is actually practiced, shows FTR still gives the lowest forgetting and highest accuracy of the four methods, with an order of magnitude less seed-to-seed variance than the next best. Together, these results invert the original study’s practical guidance: architecture-specific tuning of the stability budget is largely unnecessary, but the dual-ascent schedule that enforces it is not a free implementation detail, and checking sensitivity to it, ideally by reporting S , belongs in the standard protocol for Lagrangian-constrained continual learning.




Figure 1: Five architecturally distinct model families are usually assumed to each require their own stability budget. This paper’s central finding is that, under Lagrangian-dual constrained continual learning, the budget is instead set almost entirely by a shared control mechanism, the dual-ascent optimizer’s own schedule, common to every architecture trained with it.

1 Introduction

The fundamental challenge of continual learning is to enable neural networks to acquire new knowledge from sequential tasks without catastrophic interference with previously learned information [McCloskey and Cohen, 1989, Robins, 1995]. The stability-plasticity dilemma has motivated numerous approaches, including parameter-space regularization methods such as Elastic Weight Consolidation (EWC) [Kirkpatrick et al., 2017] and Synaptic Intelligence (SI) [Zenke et al., 2017], output-space distillation approaches like Learning without Forgetting (LwF) [Li and Hoiem, 2017], and memory-based methods including experience replay [Lopez-Paz and Ranzato, 2017] and generative replay [Shin et al., 2017].

A central practical challenge across all these methodological families is setting the magnitude of the stability constraint, typically represented as a scalar hyperparameter controlling the aggressiveness of stability protection. For EWC, practitioners must specify the Fisher coefficient. For LwF, the distillation weight must be tuned. For replay methods, the buffer size requires selection. Existing work suggests that these hyperparameters should be architecture-dependent, as loss-landscape curvature varies substantially across models of different depths, widths, and normalization strategies [Keskar et al., 2017, Ghorbani et al., 2019, Li et al., 2018].

This assumption is motivated by decades of work connecting loss-landscape geometry to optimization difficulty, generalization, and robustness [Novak et al., 2018, Li et al., 2018, Keskar et al., 2017]. The intuition is compelling: models with sharper minima (higher Hessian curvature or Fisher information) should require stronger stability pressure to prevent catastrophic forgetting because small parameter perturbations produce large functional changes. Conversely, models with flatter minima should tolerate larger updates.

Central Research Question. We ask whether this curvature-dependent hypothesis holds when stability is enforced as a hard constraint in function space via KL divergence, rather than as a soft penalty in parameter space. The constraint $D_{\text{KL}}(f_{\theta} \| f_{\text{old}}) \leq \varepsilon$ directly limits functional drift in output probability distributions, raising the question: does the critical stability budget ε^* depend on model architecture and loss-landscape geometry, or does it emerge as an invariant of the task structure alone? We additionally ask a question the constrained-continual-learning literature has not, to our knowledge, posed: *is ε^* in fact determined chiefly by the constrained optimizer’s own dual-ascent schedule rather than by the model or the task at all?* Section 5 answers this second question first, because the answer determines how the rest of the paper’s architecture-independence result should be read: not as a hyperparameter-free constant, but as a value conditional on the schedule, whose *architecture-independence* (holding that schedule fixed) is nonetheless real and is the paper’s remaining claim.

Primary Findings. Contrary to the reading that stability crossovers reflect task or architecture geometry alone, we find that the dual-ascent optimizer’s own schedule dominates: holding architecture and task fixed, a five-fold change in the dual learning rate η_{λ} , the initial multiplier λ_{init} , or the per-task step budget each moves ε^* by more than an order of magnitude (Section 5), more than the 24–75-fold range in architecture size moves it, and more than any single robustness axis we test afterward (KL direction, task ordering, class-incremental vs. task-incremental protocol). This three-dimensional schedule sensitivity is not irreducible, however: Section 13 shows it collapses substantially onto a single ratio, $S = \lambda_{\text{init}}/(\eta_{\lambda}N)$, holding which fixed keeps ε^* within a factor of 2–3 across architectures and schedule routes, versus the order-of-magnitude range any of its three components produces alone. Holding the schedule fixed, ε^* remains architecture-independent in the sense that matters operationally: no per-architecture retuning is needed, across an expanded, 30-model, five-family zoo, though a hierarchical model shows genuine (if modest) between-architecture structure that a single reported constant obscures (Section 7); notably, applying our own optimizer-schedule diagnostic to our own zoo design caught a second, self-inflicted instance of the same confound inflating that apparent structure (Section 14.3).

Contributions. This work makes the following contributions:

1. We identify and directly test a confound in Lagrangian-dual constrained continual learning that, to our knowledge, no prior work on this optimization paradigm has isolated: because the dual variable is reset at every task boundary, the apparent stability threshold of a KL-constrained method is jointly determined by the optimizer’s own schedule and by task/architecture geometry, and the schedule’s effect can dominate (Section 5).
2. Holding the schedule fixed, we expand the architecture zoo from 8 to 30 models across five representational families (CNN, ResNet, plain MLP, ViT, MLP-Mixer) and show, via a hierarchical partial-pooling model, that ε^* is not a single universal constant as v1 claimed but has real, if modest, between-architecture structure; curvature-based metrics still fail to predict it, now with adequate statistical power to support that conclusion (Section 7). We further show that part of the apparent family-level structure we initially observed was itself a second, self-inflicted instance of the schedule confound in our own zoo design, caught and corrected by applying our own diagnostic standard (Section 14.3), and that LwF’s own universal crossover replicates, while EWC and SI show no reliable crossover at all, on this same expanded, matched architecture subset (Section 9). Task ordering (Section 12) and the class-incremental protocol (Section 11) are both real, non-schedule sources of variation in ε^* that v1’s single fixed ordering and task-incremental-only design could not surface.

3. We correct a methodological error in the original F-test for constancy (its code computed p using $df_2 = 1000$ while the manuscript reported $df = (7, 7)$) with a hierarchical partial-pooling model, and show the original curvature-correlation result was underpowered rather than confirmatory of the null (Section 7).
4. We identify and test the specific ratio, $S = \lambda_{\text{init}}/(\eta_{\lambda}N)$, that governs the schedule’s effect on ε^* , and show that holding it fixed while independently varying its three components by up to $5\times$ keeps ε^* within a factor of 2–3 across five architectures spanning four representational families, converting an otherwise intractable three-hyperparameter sensitivity into a single reportable quantity (Section 13).

2 Related Work

2.1 Continual Learning Regularization

Elastic Weight Consolidation (EWC) [Kirkpatrick et al., 2017] imposes penalties on parameter changes weighted by the diagonal Fisher information matrix, thereby protecting parameters deemed important for prior tasks. Synaptic Intelligence (SI) [Zenke et al., 2017] accumulates importance weights online through path integrals over training trajectories. Both methods enforce stability in parameter space, implicitly assuming that maintaining proximity in parameter space ensures proximity in function space. This assumption fails in overparameterized neural networks, where distinct parameter vectors can produce functionally identical mappings [Garipov et al., 2018, Frankle et al., 2020, Frankle and Carbin, 2019]. Recent work on mode connectivity [Garipov et al., 2018, Draxler et al., 2018] has established the non-convexity of the loss landscape in parameter space, further motivating function-space approaches.

2.2 Function-Space and Distillation-Based Continual Learning

Learning without Forgetting (LwF) [Li and Hoiem, 2017] constrains output distributions through a KL-divergence distillation loss, maintaining the student model’s outputs close to the teacher model’s predictions. Dark Experience Replay (DER) [Buzzega et al., 2020] extends this by combining distillation with stored logits. Titsias et al. [2020] regularize continual learning directly in function space via a KL term derived from a Gaussian-process view of the network’s outputs, and Pan et al. [2020] (FROMP) regularize network outputs at a set of “memorable past” points using a GP-derived functional prior with an explicit memory. Both use a *soft*, memory-based functional prior; FTR differs in enforcing a *hard* constraint via Lagrangian dual ascent with no stored exemplars, and, unlike either, is evaluated here specifically for whether its critical budget is architecture-independent.

2.3 Constrained Learning via Lagrangian Duality

Beyond continual learning specifically, Chamon et al. [2023] formalize constrained empirical risk minimization as a general learning paradigm and prove near-PAC learnability together with empirical duality-gap bounds for non-convex losses, the theoretical setting that FTR’s Lagrangian dual-ascent update (Eqs. 2–5) implicitly relies on. Elenter et al. [2024] give near-optimal solutions for this class of constrained-learning problems, and Hounie et al. [2023] study *resilient* constrained learning, where the constraint level itself is relaxed adaptively based on cost, directly relevant to how ε is chosen here and a natural extension we do not pursue. Closest to the present work, Elenter et al. [2023] formulate continual learning explicitly as a Lagrangian-dual constrained-optimization problem (Primal-Dual Continual Learning, PDCL) with a dual variable that adapts via projected ascent

on constraint violation, evaluated on Tiny-ImageNet, SpeechCommands, and MedMNIST-OrganA in the class-incremental setting. PDCL already demonstrates that Lagrangian-dual constrained continual learning is practical and studies its dual dynamics; it does not, however, constrain drift in *function space* via a KL ball on output distributions, and does not study whether its own critical constraint level is architecture-independent, the two questions this paper is organized around.

2.4 Trust Regions and KL Constraints

Trust Region Policy Optimization (TRPO) [Schulman et al., 2015] constrains policy updates in reinforcement learning through explicit KL-divergence trust regions, motivating the FTR formulation. The connection between KL constraints and projected gradient descent in Riemannian function space has been explored in online convex optimization [Hazan, 2019, Rakhlin et al., 2012, Shalev-Shwartz, 2012]. Recent work on natural gradient methods [Amari, 1998, Martens and Grosse, 2015] establishes connections between KL constraints and information geometry on parameter manifolds.

2.5 Curvature and Forgetting: an Active, More Recent Literature

Our claim that curvature fails to predict ε^* under a function-space hard constraint should not be read against the older curvature literature alone [Keskar et al., 2017, Ghorbani et al., 2019, Li et al., 2018]; a more recent, active line of work directly studies curvature-forgetting links that our result must be positioned against rather than ignore. Bian et al. [2024] use sharpness-aware minimization to reduce forgetting by flattening the loss landscape, and its extension pushes this further; other work explores Hessian eigenvalue spectra as a forgetting diagnostic, and Deng et al. [2021] project gradients using curvature information to reduce interference. Most directly relevant, Lanzillotta et al. [2023] formally derive that sharper old-task minima require a *tighter* parameter-space constraint to prevent forgetting. That result is not in tension with the present finding: it is a parameter-space claim, and our finding that curvature fails to predict ε^* is specific to *function-space* hard constraints. Lewandowski et al. [2025] and Lewandowski et al. [2023] study curvature’s relationship to loss of plasticity (the stability-plasticity tradeoff’s other side), a complementary rather than contradictory line.

2.6 Phase Transitions in Deep Learning

Phase transitions in neural network training have been documented in contexts including grokking [Power et al., 2022], lottery-ticket emergence [Frankle and Carbin, 2019], and loss-landscape structure [Li et al., 2018, Keskar et al., 2017]. Recent finite-size-scaling treatments of grokking [Wang, 2026] and of absorbing phase transitions in deep networks [Tamai et al., 2025] apply statistical-mechanics tools, tracking how a transition sharpens with system size, directly analogous to the width-family analysis in Section 8. A thematically adjacent but mechanistically different framing casts the stability-plasticity tradeoff itself as a Stefan (moving-boundary) problem governed by a single scalar [Khilar, 2026]; we note it here as a kindred "physics-of-training" perspective rather than a directly competing method.

3 Method: Functional Trust Regions

3.1 Problem Formulation

Consider a sequence of tasks $\{T_1, T_2, \dots, T_m\}$ presented sequentially. At task $t + 1$, the learner has access to training data from the current task but not from prior tasks. The parametric model is

$f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$, where $\theta \in \mathbb{R}^d$ are model parameters. Let f_{old} denote the model trained on tasks $\{1, \dots, t\}$.

Task-incremental vs. class-incremental. v1 did not state this explicitly. All main results use the *task-incremental* protocol: the output head is sized to the number of classes *within* a single task (e.g. 2 for the binary CIFAR-10 splits), labels are remapped locally per task (class c of task t becomes local label $0, 1, \dots$), and evaluation at test time uses an oracle task ID to select both the correct label remapping and, implicitly, which classes are eligible outputs. This is a materially easier protocol than class-incremental learning, where a single shared head must discriminate among all classes seen so far with no oracle task ID at test time. Section 11 reports a class-incremental variant on a subset of architectures to test whether the crossover survives the harder protocol.

3.2 Constrained Optimization Formulation

We formulate task $t + 1$ learning as a constrained optimization problem:

$$\theta_{t+1} = \arg \min_{\theta} \mathcal{L}_{\text{task}}(\theta) \quad \text{subject to} \quad D_{\text{KL}}(f_\theta \| f_{\text{old}}) \leq \varepsilon, \quad (1)$$

where $D_{\text{KL}}(f_\theta \| f_{\text{old}}) = \mathbb{E}_{x \sim \mathcal{D}_t} [D_{\text{KL}}(f_\theta(x) \| f_{\text{old}}(x))]$ is the expected KL divergence between output distributions, estimated on the current task’s training batch. The constraint radius ε represents the stability budget in function space.

3.3 Optimization via Lagrangian Dual Ascent

We solve the constrained problem using Lagrangian relaxation. The Lagrangian is:

$$\mathcal{L}(\theta, \lambda) = \mathcal{L}_{\text{task}}(\theta) + \lambda(D_{\text{KL}}(f_\theta \| f_{\text{old}}) - \varepsilon), \quad (2)$$

where $\lambda \geq 0$ is the dual variable. We employ dual gradient ascent with momentum:

$$\theta_{t+1} \leftarrow \theta_t - \eta_\theta \nabla_\theta \mathcal{L}(\theta_t, \lambda_t), \quad (3)$$

$$m_{t+1} \leftarrow \rho m_t + \eta_\lambda (D_{\text{KL}} - \varepsilon), \quad (4)$$

$$\lambda_{t+1} \leftarrow \max(0, \lambda_t + m_{t+1}), \quad (5)$$

where η_θ is the primal learning rate, η_λ the dual learning rate, and ρ the momentum coefficient. The dual variable adapts automatically: when drift exceeds ε , λ increases (tightening the constraint); when the model is stable, λ decreases (relaxing it).

λ is reset at every task boundary. A detail omitted from v1’s Method section but essential to the dynamics analyzed in Section 5: λ is re-initialized to λ_{init} at the start of *every* task transition, not held as a single running state across the whole training sequence. Within a task, each step compares the current drift D_{KL} against ε and nudges λ by η_λ times an EMA of the violation $(D_{\text{KL}} - \varepsilon)$ (Eq. 5). Two regimes follow directly from this:

- **Tight ε :** the drift an *unconstrained* model would accumulate while fitting a new task typically exceeds ε within the first few post-warmup steps, so the violation is positive, λ is driven upward (up to λ_{max}), and the constraint stays actively engaged for the remainder of the task, so forgetting stays low.

- **Loose ε :** drift stays under ε throughout, the violation is persistently negative, and λ decays toward 0 within a number of steps set by λ_{init} , η_λ , and the momentum ρ ; once $\lambda \approx 0$ the constraint is effectively disengaged and training proceeds as unconstrained fine-tuning for whatever fraction of the task’s step budget remains.

The crossover ε^* is, mechanically, the value of ε at which a task transitions from the first regime to the second within the fixed per-task step budget. This makes ε^* a joint function of (a) how much functional drift unconstrained gradient descent on the new task would produce, which is a property of the task and architecture, and (b) how many steps it takes λ to decay from λ_{init} to 0 under slack, which is a property of η_λ , λ_{init} , λ_{max} , ρ , the per-task step budget, and the temperature T (which rescales D_{KL} itself, since the drift term used in training is $T^2 \cdot D_{\text{KL}}$). Section 5 tests directly whether (b) has a material effect on ε^* , which v1 did not check.




Figure 2: Stylized illustration of the mechanism above (not measured data: a toy simulation of Eqs. 2–5 with an illustrative, monotonically decreasing drift-rate function $\bar{d}(\lambda)$, used only to make the two regimes visually concrete before the measured results in Section 5). (A) λ ’s trajectory within a single task for five values of ε : tight budgets keep λ engaged near λ_{init} or rising; loose budgets drive it to 0 before the task’s step budget N is exhausted. (B) The resulting crossover in λ at the end of the task as a function of ε , for this toy simulation.

3.4 Relationship to Learning without Forgetting

When λ is held fixed rather than adapted, FTR reduces exactly to Learning without Forgetting, which minimizes:

$$\mathcal{L}_{\text{LwF}}(\theta) = \mathcal{L}_{\text{task}}(\theta) + \lambda D_{\text{KL}}(f_\theta \| f_{\text{old}}). \quad (6)$$

The key distinction is that FTR enforces a hard constraint with adaptive dual coefficient, providing an interpretable stability budget ε that practitioners can reason about directly, whereas LwF requires grid search over fixed distillation coefficients.

3.5 Implementation Details

All experiments employ the following hyperparameters unless stated otherwise: $\lambda_{\text{init}} = 1.0$, $\lambda_{\text{max}} = 50.0$, $\eta_\lambda = 0.005$, momentum $\rho = 0.9$, softmax temperature $T = 2.0$, and one epoch of warmup per task. The task loss $\mathcal{L}_{\text{task}}$ uses cross-entropy with Adam [Kingma and Ba, 2015], learning rate

$\eta_\theta = 0.001$. Section 5 sweeps every one of these choices individually against a dense ε grid to test their effect on ε^* .

4 Experimental Setup

4.1 Architecture Zoo

v1’s zoo comprised eight architectures, all convolutional or residual. We expand it to 30 architectures spanning five representational families and a $75.8\times$ parameter range (36,866–2,795,554): a CNN width sweep (W8 through W128, 8 architectures), a CNN depth sweep (2–5 conv layers at width 32, 4 architectures), CNN variants without BatchNorm at three scales, a ResNet-18-style family with 2 residual blocks/layer at widths 8/16/24/32, a ResNet-Lite family with 1 block/layer at the same four widths (added specifically to fill the curvature gap between the CNN family and the original single ResNet-18 outlier, discussed below) including one no-BatchNorm variant, a plain MLP with no spatial inductive bias at two hidden widths, and two scales each of a patch-based Vision Transformer and an MLP-Mixer. Full per-architecture parameter counts, Hessian trace, Fisher trace, spectral norm, and d_{eff} are in Table 10 (Appendix A). Hessian trace still spans a comparable range to v1 ($39\times$: 79.7 at CNN_W32_NoBN to 3110 at ResNet18_W8) but is now populated by all 30 architectures rather than 8, and by five families with genuinely different inductive biases rather than two families that are architecturally close relatives.

4.2 Dataset and Task Construction

All main experiments use CIFAR-10 [Krizhevsky, 2009] partitioned into five sequential binary classification tasks, where each task contains two classes. Each task contains 1,000 samples per class (2,000 total), trained for 4 epochs/task for the CNN, ResNet, ResNet-Lite, and MLP families, and 5 epochs/task for ViT and MLP-Mixer (a per-family accommodation made when the zoo was built, whose consequence for ε^* we return to in Section 14.3). The forgetting metric is computed as the mean decrease in accuracy on all previously learned tasks after completion of training on the final task; we additionally report new-task accuracy (accuracy on task t immediately after training on task t , before any subsequent forgetting) to rule out the trivial explanation that low forgetting at small ε reflects a model that has simply failed to learn the new task (Sec. 7.5).

Four further robustness checks (a CIFAR-100 classes-per-task sweep, multiple random class-to-task orderings, a KL-direction ablation, and a class-incremental variant with a single shared head) are reported in Sections 12, 10, and 11 once the main results (Sections 5–8) have established the vocabulary needed to interpret them.

5 Diagnostic: Is ε^* a Task-Structure Invariant or an Optimizer-Schedule Artifact?

Section 3 showed that λ resets to λ_{init} at every task boundary and then rises or decays within the task depending on whether measured drift exceeds ε . This makes ε^* , mechanically, a function of *two* things at once: how much functional drift unconstrained learning on a new task would produce (task/architecture geometry), and how many steps it takes λ to travel from λ_{init} to 0 (or to λ_{max}) under the dual-ascent rule (an optimizer-schedule property: η_λ , λ_{init} , λ_{max} , momentum ρ , steps/task, and the temperature T , which rescales the drift term itself). v1 varied only the first axis. We test the second directly.

5.1 Setup

On two anchor architectures (CNN_W16, CNN_W32) and a 10-point ε grid $\{1, 3, 5, 6, 7, 8, 9, 10, 12, 20\}$, 3 seeds, we ran the **baseline** configuration of Section 3 ($\eta_\lambda = 0.005$, $\lambda_{\text{init}} = 1.0$, $\lambda_{\text{max}} = 50$, $\rho = 0.9$, $T = 2.0$, 4 epochs/task) alongside eleven single-factor perturbations: η_λ and λ_{init} each $\times 5$ and $\div 5$; $\lambda_{\text{max}} \in \{10, 200\}$; epochs/task $\in \{2, 8\}$ (baseline 4); $\rho = 0.5$; and $T \in \{1.0, 4.0\}$. For each of the 24 (architecture, condition) cells we fit the same logistic-sigmoid procedure as Section 6 and bootstrap a 95% CI (1,000 resamples over seeds). Four conditions produced a forgetting curve that was *flat across the entire* $[1, 20]$ grid (range < 0.05) rather than an S-shape inside it, so we reran those four plus baseline on a wider, coarser grid extending to $\varepsilon = 100$ ($\{0.05, 0.1, 0.3, 0.5, 1, 3, 5, 7, 10, 15, 20, 30, 40, 50, 75, 100\}$, same architectures and seeds) to locate the true crossover rather than report a degenerate sigmoid fit.

5.2 Results

Table 1: ε^* under single-factor perturbations of the dual-ascent schedule, holding architecture and task fixed (CIFAR-10, 2 anchor architectures, 3 seeds, 95% bootstrap CI). “Ratio” is ε^* relative to that architecture’s own baseline. On the original dense grid ($\varepsilon \in [1, 20]$), five rows produced a forgetting curve flat across the *entire* grid rather than an S-shape inside it, so no transition was located; we rescanned exactly those five (condition, architecture) cells with a wider grid ($\varepsilon \in [0.05, 100]$, marked ‡) to locate the true crossover rather than report only a direction.

Condition	ε^* (CNN_W16)	R^2	ratio	ε^* (CNN_W32)	R^2	ratio
baseline	13.09 [11.2, 54.4]	0.995	1.00	12.45 [10.9, 54.4]	0.991	1.00
$\eta_\lambda \times 5$	2.70 [2.7, 3.2]	0.970	0.21	0.97 [0.7, 1.3]	0.955	0.08
$\eta_\lambda \div 5$	79.5 [61.4, 147.6] ‡	0.993	6.07	58.5 [55.7, 71.5] ‡	0.991	4.70
$\lambda_{\text{init}} \times 5$	76.7 [69.3, 85.0] ‡	0.915	5.86	90.1 [67.7, 158.4] ‡	0.875	7.24
$\lambda_{\text{init}} \div 5$	1.69 [1.05, 2.10] ‡	0.970	0.13	1.04 [1.00, 1.83] ‡	0.951	0.08
$\lambda_{\text{max}} = 10$	13.69 [11.5, 54.4]	0.986	1.05	15.28 [11.2, 54.4]	0.984	1.23
$\lambda_{\text{max}} = 200$	13.06 [11.9, 14.2]	0.989	1.00	15.55 [11.3, 54.4]	0.990	1.25
epochs/task = 8 ($\times 2$)	3.25 [3.0, 3.5]	0.992	0.25	3.66 [3.2, 4.4]	0.980	0.29
epochs/task = 2 ($\div 2$)	16.3 [13.3, 38.1] ‡	0.872	1.25	36.0 [0.8, 46.8] ‡ ^a	0.305	2.89
momentum $\rho = 0.5$	8.06 [7.6, 10.8]	0.994	0.62	6.97 [6.6, 7.5]	0.974	0.56
$T = 4.0$ ($\times 2$)	15.56 [11.6, 33.5]	0.991	1.19	15.32 [12.2, 21.9]	0.996	1.23
$T = 1.0$ ($\div 2$)	13.51 [10.4, 19.4]	0.990	1.03	12.13 [10.6, 22.4]	0.970	0.97

^aSigmoid $R^2 = 0.31$ even on the wide grid for this one cell (CNN_W32, epochs/task=2); the point estimate is a rough indication of direction and magnitude, not a tightly located transition.

Three of the six hyperparameters we varied, η_λ , λ_{init} , and epochs/task, move ε^* far more than architecture does in this paper’s own main result, and with the wide-grid follow-up all three directions are now precisely located rather than bounded qualitatively. Doubling epochs/task or raising η_λ five-fold collapses ε^* from ≈ 13 to 2–4 (a 4–13 $\times$ reduction, tight bootstrap CIs, sigmoid $R^2 \geq 0.97$). Halving epochs/task, lowering η_λ five-fold, or raising λ_{init} five-fold each push ε^* up by 1.25–7.2 $\times$, as high as ≈ 90 ; lowering λ_{init} five-fold pushes it down by 7–12 $\times$, to ≈ 1 –1.7. Taken together, across the six single-factor perturbations tested, ε^* spans roughly 1.0 to 90 on the *same* architecture and task: essentially two orders of magnitude from optimizer hyperparameters alone. Momentum ρ has a smaller but still real effect (≈ 0.56 – $0.62\times$). Only λ_{max} and the softmax temperature T are inert (ratios within $[0.97, 1.25]$, CIs overlapping baseline): λ_{max} because the dual variable rarely approaches its cap under these dynamics, and T because it rescales both the drift term used in training *and* the quantity compared against ε , so its effect cancels at the crossover.




Figure 3: ε^* (log scale, 95% bootstrap CI) under each single-factor schedule perturbation, CNN_W16 and CNN_W32, sorted by CNN_W16’s value. Dashed line marks the CNN_W16 baseline. $\lambda_{\max}$ and T rows cluster tightly around baseline; every other row is displaced by at least one order of magnitude on the log axis.

5.3 Generalizing Beyond CNNs

Everything above uses two CNN width variants as anchors. Before treating the schedule confound as a general property of Lagrangian-dual constrained continual learning rather than a CNN-specific quirk, we repeat the two cleanest single-factor perturbations, $\eta_\lambda \times 5$ and epochs/task $\times 2$, on one anchor from each of the four remaining families (ResNet18_W16, MLP_H128, ViT_Tiny, Mixer_Tiny), 3 seeds, dense ε grid, epochs/task held at 4 for all four (so this test is itself run on a schedule-matched, not confounded, sub-zoo).

Both perturbations collapse ε^* by roughly an order of magnitude on every family, and to a strikingly similar *absolute* value regardless of starting architecture: $\eta_\lambda \times 5$ gives $\varepsilon^* = 2.23$ (ResNet18_W16), 3.15 (MLP_H128), 2.97 (ViT_Tiny), 2.82 (Mixer_Tiny), all within $[2.2, 3.2]$, $R^2 \geq 0.88$ throughout. Doubling epochs/task gives 3.04, 3.40, 2.44, 3.15 respectively, all within $[2.4, 3.4]$, $R^2 \geq 0.92$ throughout. This is a tighter cluster across four architecturally distinct families (convolutional, fully-connected, attention-based, token-mixing) than the same two conditions produced across the two CNN widths in Table 1 (2.70/0.97 for $\eta_\lambda \times 5$; 3.25/3.66 for epochs $\times 2$), so the schedule confound is, if anything, more robust across architecture families than within a single one. The remaining two conditions, $\lambda_{\text{init}} \times 5$ and $\div 5$, and the unperturbed baseline, are less reliably estimated on this family sub-zoo: three of the sixteen (family, condition) cells here hit the same curve_fit bound-saturation issue diagnosed in Section 6 (MLP_H128 and ViT_Tiny baselines, both

pinned at the grid’s upper bound) or land at poor fit quality ($R^2 < 0.5$ for ResNet18_W16 and ViT_Tiny under $\lambda_{\text{init}} \times 5$, which were not rescanned on the wider grid Section 5’s CNN-only version used for its own analogous cells). We do not draw quantitative conclusions from those specific cells, but every one that fits reliably moves in the predicted direction, and the two best-established conditions generalize cleanly across all five representational families in the zoo.

5.4 Conclusion

This reframes, rather than negates, the claim that ε^* reflects task and architecture. Section 6 still asks and answers a well-posed narrower question: holding the optimizer schedule fixed, is ε^* architecture-independent? Per Section 7, the answer on the expanded architecture zoo is still substantially yes. What changes is the status of the numeric value itself. At fixed architecture and fixed task, changing only how fast the dual variable is allowed to move, a choice most papers using Lagrangian-relaxed constraints (including v1 of this one) report as a fixed implementation detail rather than a variable to sweep, shifts the critical stability budget by more than an order of magnitude, comfortably exceeding the effect that a 24–75 $\times$ range in architecture size produced in Section 6. So $\varepsilon^* \approx 7\text{--}13$ is not a hyperparameter-free constant of CIFAR-10’s task structure that a practitioner can adopt off the shelf; it is the value produced by one particular, previously unremarked choice of η_λ , λ_{init} , and steps per task, and a different reasonable choice of any of those three moves it by more than an order of magnitude.

This confound is specific to methods whose stability coefficient is an *adaptive dual variable reset at each task boundary*. It does not apply to LwF, whose distillation weight is a fixed, grid-searched hyperparameter with no reset dynamics to confound, or to EWC and SI, which show no crossover at all (Section 9) and so raise no analogous question. It applies directly to FTR and, by the same argument, to PDCL [Elenter et al., 2023] and any other Lagrangian-dual constrained-learning method with a periodically reset dual variable. To our knowledge, no prior work on Lagrangian-dual constrained continual learning (Elenter et al., 2023; Chamon et al., 2023; Elenter et al., 2024, Hounie et al., 2023) has isolated the dual-ascent schedule itself as a confound for an apparent stability threshold, and we are not aware of this question being asked elsewhere in the constrained-learning literature.

The practical upshot inverts v1’s: architecture-specific tuning of ε is indeed unnecessary, but optimizer-schedule-specific tuning is not, and is currently undocumented practice across the distillation-based continual-learning literature that uses dual ascent or an equivalent adaptive coefficient. Section 13 shows this need not be a counsel of despair: the schedule’s three-dimensional effect on ε^* collapses substantially, though not completely, onto a single reportable ratio, $S = \lambda_{\text{init}}/(\eta_\lambda N)$, which, held fixed, keeps ε^* within roughly a factor of 2–3 across architectures and schedule routes rather than the order-of-magnitude range any one of its three components produces alone.

6 The Stability Crossover

Holding the dual-ascent schedule fixed at the Section 3 defaults (the same schedule used throughout v1), we ran the dense 17-point ε sweep $\{0.5, 1, 2, 3, 4, 5, 6, 6.5, 7, 7.5, 8, 9, 10, 12, 15, 25, 50\}$ on all 30 architectures of the expanded zoo (Section 4), 5 seeds per architecture (10 for the two diagnostic anchor architectures), fit the same logistic-sigmoid procedure as v1 (Eq. 7 of the original manuscript, unchanged), and bootstrapped a 95% CI over seeds (1,000 resamples). Every architecture produces a clean sigmoid ($R^2 \geq 0.78$, 27/30 with $R^2 \geq 0.94$); Figure 4 plots all 30 curves and Figure 5 plots ε^* with its 95% CI for every architecture, colored by family. Table 2 reports ε^* , its bootstrap CI,

and the sigmoid sharpness k for a representative subset spanning every family; the full 30-row table is in Appendix A.




Figure 4: Dense ε sweep forgetting curves for all 30 architectures, colored by family. CNN/ResNet/MLP curves (blue/orange/green) cluster around a shared crossover region; ViT (red) and MLP-Mixer (purple) transition at visibly lower ε .

Two architectures, ResNet18_W16 and ResNetLite_W8_NoBN, have bootstrap CI upper bounds that land exactly on $50e = 135.91$, the `curve_fit` search-bound value diagnosed in Section 5; here the *point estimate* is not pinned (both are ordinary interior fits, $R^2 = 0.94$ and 0.78), but enough individual bootstrap resamples produce a flat-enough curve to hit the bound that the reported upper CI is not trustworthy. We reran both on a wider grid extending to $\varepsilon = 200$ (5 seeds, 1,000 bootstrap resamples): ResNet18_W16 relocates to $\varepsilon^* = 9.22$ [7.59, 10.67] ($R^2 = 0.953$), a much tighter interval and a point-estimate shift of -5.37 from the original narrow-grid value; ResNetLite_W8_NoBN relocates to 12.50 [6.87, 97.02] ($R^2 = 0.889$), a shift of -3.47 , with the CI substantially narrower than before but still wide, indicating this architecture’s crossover remains less precisely determined than most of the zoo even with the wider grid. We use these relocated values in place of the narrow-grid ones throughout Section 7.

Across all 30 architectures as originally recorded (narrow grid, before either the epoch-matching or wide-grid corrections), ε^* ranges [6.10, 21.45], mean 11.14, raw CV 26.8%, more than five times v1’s reported 4.96% on eight architectures. This is not primarily a statistical-power artifact of the smaller original sample: the expanded zoo includes two representational families (ViT, MLP-Mixer) with no analogue in v1’s CNN/ResNet-only zoo, and Section 7 shows the extra dispersion tracks family identity, not any tested curvature statistic, though, as that section also shows, part of it tracks two rounds of our own schedule confound rather than architecture identity itself.




Figure 5: ε^* with 95% bootstrap CI for all 30 architectures, grouped and colored by family, sorted within family by value. ViT and MLP-Mixer form a visually distinct low cluster; CNN, ResNet(+Lite), and MLP overlap in a broader band. CNN_W8_NoBN and two ResNet points show markedly wider CIs.

7 Universality Analysis

7.1 Family Structure: Mostly a Second Instance of the Diagnostic’s Own Finding

Grouping ε^* by architecture family as originally recorded (Table 3) shows the CNN, ResNet, and MLP families clustering in a broad 8–15 band, while ViT and MLP-Mixer sit systematically lower, 6–7.4, with almost no within-family spread ($\text{std} \leq 0.05$ for both). **This clustering is substantially an artifact of our own experimental design, not primarily a representational-paradigm effect:** as Section 14.3 shows, ViT and MLP-Mixer were assigned 5 epochs/task rather than the 4 used elsewhere in the zoo, and Equation 9 predicts, and a direct epoch-matched re-run confirms, that this alone accounts for most of the apparent gap: epoch-matched, ViT_Tiny and Mixer_Tiny move from 7.31/6.10 to 10.18/10.41, landing inside the CNN/ResNet/MLP family’s range rather than below it. We caught this only by applying Section 5’s own standard to our own zoo, and report it plainly rather than let the originally-recorded clustering stand as an architectural finding. One architecture, CNN_W8_NoBN, remains a clear outlier at $\varepsilon^* = 21.45$, the smallest, narrowest, unnormalized network in the zoo, which Section 7.3 shows has outsized leverage on the raw CV but does not materially move the hierarchical model’s population estimate.

Table 2: Sigmoid-fitted ε^* across a representative subset of the 30-architecture zoo (full table in Appendix A). Families: CNN (width/depth/BatchNorm sweeps), ResNet (2 blocks/layer), ResNet-Lite (1 block/layer), MLP (no spatial inductive bias), ViT (patch attention), Mixer (token/channel MLP-mixing).

Architecture	Family	Params	ε^*	95% CI	R^2	k
CNN_W8_NoBN	cnn	36,866	21.45	[18.57, 30.88]	0.986	2.35
CNN_W16	cnn	80,418	10.80	[10.03, 11.48]	0.990	4.04
Mixer_Tiny	mixer	45,442	6.10	[5.88, 6.50]	0.972	4.04
ViT_Tiny	vit	74,562	7.31	[6.35, 8.03]	0.974	4.19
ResNetLite_W8	resnet	77,482	10.46	[8.18, 13.41]	0.962	3.75
ResNet18_W8	resnet	175,882	13.76	[10.98, 26.73]	0.901	2.16
MLP_H128	mlp	410,114	14.56	[13.74, 15.37]	0.994	2.10
ResNet18_W16	resnet	700,434	14.59	[9.65, 135.91] [†]	0.937	1.63
CNN_W128	cnn	1,414,786	10.43	[9.85, 10.94]	0.979	3.84
ResNet18_W32	resnet	2,795,554	8.88	[7.45, 11.84]	0.947	2.33

[†]Bootstrap CI pinned against the `curve_fit` search bound ($50e = 135.91$) in its upper tail; see the wide-grid relocation below.

Table 3: ε^* by architecture family, full 30-architecture zoo, as originally recorded (5 epochs/task for ViT/Mixer, 4 for every other family; see Section 14.3 and Table 8 for why this comparison is confounded and what changes when it is fixed).

Family	n	mean ε^*	std
MLP-Mixer	2	6.11	0.00
ViT	2	7.36	0.04
ResNet (incl. Lite)	9	11.55	2.64
CNN	15	11.75	2.64
MLP	2	13.54	1.02

7.2 Hierarchical Partial-Pooling Model

v1’s F-test for constancy (its Table 3) computed $p = 1 - F_{\text{cdf}}(0.563; df_1=7, df_2=1000) = 0.786$ in code, while the manuscript text reported $df = (7, 7)$: an inconsistency between the stated methodology and the actual computation that we caught while reproducing it (with $df_2 = 7$ the same statistic gives $p = 0.767$, not 0.786; the discrepancy is not merely rounding). $df_2 = 1000$ has no stated justification and appears to be an unexplained placeholder. We replace it with a normal-normal hierarchical partial-pooling model: each architecture’s true crossover $\theta_a \sim \mathcal{N}(\mu, \tau^2)$, with the sigmoid point estimate observed as $\hat{\varepsilon}_a^* \sim \mathcal{N}(\theta_a, s_a^2)$ where s_a is its bootstrap standard error (fit by Gibbs sampling, 30,000 post-burn-in draws).

Fit on the 30-architecture zoo as originally recorded (i.e. before either correction below), this model gives $\mu = 10.46 \pm 0.45$ (95% CI [9.58, 11.36]), $\tau = 2.13$ (95% CI [1.56, 2.92]), and an ICC-style universality index $\tau^2/(\tau^2 + \bar{s}^2) = 0.041$ (95% CI [0.022, 0.073]): τ ’s credible interval excludes zero, suggesting real between-architecture structure. But this zoo contains two known, independently-diagnosed data-quality issues that we do not treat as acceptable to leave uncorrected before reporting a final estimate: the epoch/task schedule confound on four architectures (Section 14.3), and two architectures whose bootstrap CI pins against the `curve_fit` search bound (Section 6). Substituting the epoch-matched values for ViT_Tiny, Mixer_Tiny, ViT_Small, and Mixer_Small (Table 8) and

Table 4: ε^* by architecture family, corrected: ViT_Tiny, Mixer_Tiny, ViT_Small, and Mixer_Small replaced with their epoch-matched values (Table 8); ResNet18_W16 and ResNetLite_W8_NoBN replaced with their wide-grid-relocated values (Section 6); CNN and MLP families unchanged. All five families now overlap in a single 10.5–13.5 band.

Family	n	mean ε^*	std
MLP-Mixer	2	10.52	0.11
ResNet (incl. Lite)	9	10.57	1.87
ViT	2	11.30	1.13
CNN	15	11.75	2.64
MLP	2	13.54	1.02

Family clustering before and after correcting the two schedule/CI data-quality issues (Sections 13.2 and 6)




Figure 6: ε^* by architecture and family, as originally recorded (left) versus corrected (right); vertical bars mark family means. ViT and MLP-Mixer move from a visually distinct low cluster into the same band as the rest of the zoo; CNN and MLP are unchanged. CNN_W8_NoBN (isolated point in the CNN row) remains an outlier in both panels.

the wide-grid-relocated values for ResNet18_W16 and ResNetLite_W8_NoBN, then refitting the identical model on the same 30 architectures, gives $\mu = 11.07 \pm 0.25$ (95% CI [10.58, 11.55]), $\tau = 1.03$ (95% CI [0.68, 1.50]), and ICC = 0.052 (95% CI [0.023, 0.102]).

Correcting the schedule confound cuts τ by more than half, from 2.13 to 1.03: most of the raw between-architecture spread the uncorrected zoo displayed was schedule-driven, the same lesson as Section 5’s original diagnostic recurring one level up. The ICC, however, does *not* fall by a comparable factor, and in fact sits slightly *above* the uncorrected estimate (0.052 vs. 0.041). This is not a contradiction: $\text{ICC} = \tau^2 / (\tau^2 + \bar{s}^2)$ depends on the average *measurement* noise $\bar{s}^2$ as much as on τ , and the wide-grid correction sharply reduces $\bar{s}^2$: ResNet18_W16’s bootstrap SE drops from 36.6 to 0.75 and ResNetLite_W8_NoBN’s from 42.6 to 22.3, both far more precise than their original, bound-saturated estimates. A more precisely measured zoo is a more sensitive test, and it detects the same, now smaller, τ against a quieter noise floor. Both corrected quantities are consistent with the same qualitative reading: the credible interval for τ still excludes zero, so we do not claim ε^* is a literal universal constant, but $\tau \approx 1$ on a population with mean ≈ 11 is a small

effect in absolute terms (all five family means in Table 4 fall within a 10.5–13.5 band), and both the ICC before and after correction round to "small" on standard interpretive scales (< 0.1). We report both the raw and corrected fits rather than only the more favorable one; the corrected values are used in Sections 7.3 onward unless stated otherwise.

7.3 Leave-One-Out Sensitivity

On the zoo as originally recorded, dropping each architecture in turn and recomputing the raw CV shows two overlapping sources of leverage: the single largest is CNN_W8_NoBN (CV drops from 26.8% to 21.6%, -5.2 points), consistent with it being a genuine, if extreme, outlier; the next two are Mixer_Tiny and Mixer_Small (-1.3 points each), i.e. the same two confounded architectures Section 14.3 identifies, which were pulling the mean down and therefore also carried above-average leverage on the raw CV.

Repeating this on the fully-corrected data (epoch-matching and wide-grid relocation both applied, Table 4), the corrected CV is 20.2%, and CNN_W8_NoBN remains the single largest-leverage point by a wide margin (removing it drops CV to 12.4%, -7.8 points), still consistent with a genuine, if extreme, outlier rather than a typical zoo member. Critically, the leverage previously carried by Mixer_Tiny and Mixer_Small disappears once the confound is corrected: no other single architecture moves the corrected CV by more than 0.5 points, so no single point drives the corrected result, unlike v1's original zoo, where ResNet18_W8 alone carried most of the curvature range. (On the original 8-architecture data, dropping ResNet18_W8 moves CV from 4.96% to 3.97%, confirming that result was not solely an artifact of that one leverage point either.)

7.4 Curvature Correlations: Adequately Powered This Time

v1's central "falsification" claim, that no curvature metric predicts ε^* , rested on $n = 8$ architectures, where the minimum $|r|$ detectable at 80% power is 0.849; its own best candidate (gradient norm, $r = 0.687$, $p = 0.060$) fell just short of significance, which at that sample size is underpowered, not confirmatory of the null. Reproducing that exact test on the original 8-architecture data with Bayes factors (JZS prior) makes this explicit: gradient norm gives $\text{BF}_{10} = 1.96$, Hessian trace 1.56, Fisher trace 1.78, spectral norm 1.46, all in the "barely worth mentioning" range and, if anything, weakly *for* a relationship, not against one.

On the full 30-architecture zoo, the same test is properly powered: minimum detectable $|r| = 0.492$ at 80% power. Recomputing on the corrected ε^* values (Table 4), every correlation remains well below this detection threshold: Hessian trace $r = -0.021$ ($\text{BF}_{10} = 0.23$), Fisher trace $r = -0.037$ (0.23), spectral norm $r = 0.099$ (0.26), gradient norm $r = -0.116$ (0.27) are all small and Bayes-negative, consistent with the null. d_{eff} ($r = -0.320$, $\text{BF}_{10} = 0.93$) and parameter count ($r = -0.306$, $\text{BF}_{10} = 0.83$) are the two exceptions: still below the well-powered detection threshold and still short of conventional significance ($p = 0.085$ and 0.100), but their Bayes factors sit in the "barely worth mentioning" range rather than clearly favoring the null, and both moved further from zero once the two data-quality corrections were applied (from $r = -0.091$ and -0.137 on the uncorrected zoo). We flag this rather than round it down to a clean null: parameter count and effective dimension are the two metrics most directly tied to model *size*, and size is also what differs most systematically between the low- ε^* MLP-Mixer/ResNet cluster and the high- ε^* CNN/MLP cluster in Table 4, so a modest, sub-threshold size association is plausible without being a confirmed effect at this sample size. The four purely curvature-based metrics (Hessian trace, Fisher trace, spectral norm, gradient norm) show no such trend in either direction. (ViT's fused attention kernel does not support the double-backward the Hessian/spectral estimators require by default; we force

the math-mode SDPA backend during curvature measurement only, Appendix B, recovering all 30 architectures rather than excluding ViT.) Unlike v1’s claim, this is now an adequately powered result: loss-landscape curvature specifically does not predict ε^* on this zoo, though raw model size may carry a small, currently sub-significant signal worth testing at larger n . What curvature misses, family identity partially captures (Section 7.1): the relevant architectural property is closer to representational paradigm (convolution vs. attention vs. token-mixing) than to any scalar loss-landscape statistic tested here or in v1, though Section 14.3 shows part of that apparent family effect is itself a schedule confound.

7.5 Rule Out the Trivial Explanation: New-Task Accuracy

A low-forgetting regime is only meaningful if the model is still learning the new task, not simply failing to update at all. Across all 30 architectures, mean accuracy on the just-learned task (measured immediately after training on it, before any subsequent forgetting) is 77.8% at tight budgets ($\varepsilon \leq 3$, $n = 120$ cells) versus 80.8% at loose budgets ($\varepsilon \geq 25$, $n = 60$ cells), a 3-point difference, with the tightest individual architecture-budget cell never dropping below 67.1% new-task accuracy. The low-forgetting regime is genuine stability-*with*-plasticity, not a degenerate case of the constraint preventing learning altogether.

8 Finite-Size Scaling of the Transition

Testing whether the sigmoid sharpness k grows with network width, the finite-size-scaling signature that would indicate the crossover sharpens toward a true phase transition as width $\rightarrow \infty$, on the full 8-point CNN width family (W8 through W128, same task, same schedule, 5 seeds each): $k \sim W^\alpha$ with $\alpha = 0.124 \pm 0.082$, $R^2 = 0.276$, $p = 0.181$ (Figure 7). This is not statistically significant, and we report it as a null result rather than force a narrative: at the widths tested here, we find no reliable evidence that the FTR crossover sharpens with system size. This does not rule out finite-size scaling at larger widths than tested, but the data in hand do not support the claim.




Figure 7: Left: sigmoid sharpness k vs. CNN width on log-log axes; the fitted power law is not statistically significant. Right: ε^* vs. width for the same family, schedule fixed; consistent with the narrow band reported in Section 6.

9 Cross-Method Comparison

v1 compared FTR, LwF, and EWC on different architecture subsets (FTR on all 8, LwF/EWC on 5, omitting, among others, CNN_W32_NoBN from the LwF/EWC comparison). We rerun FTR, LwF, EWC, and SI on the identical 15-architecture subset spanning every family (Section 4, SUBSET_ARCHS), 3 seeds, dense hyperparameter sweeps for each method’s own stability knob: $\varepsilon \in [0.5, 50]$ (17 points) for FTR, distillation weight $\alpha \in [0.05, 5.0]$ (14 points) for LwF, Fisher coefficient $\lambda_{\text{EWC}} \in [1, 10^4]$ (8 points, four orders of magnitude, matching v1’s own EWC sweep range) for EWC, and SI coefficient $c \in [0.01, 50]$ (8 points) for SI.

Table 5: Cross-method comparison on the identical 15-architecture subset. Sharpness/fit-quality columns report the median sigmoid R^2 across architectures; “crossover” is each method’s own stability-knob analogue of ε^* .

Method	Mean crossover	CV	Median R^2	Interpretation
FTR (function, hard)	11.4	20.5%	≥ 0.94	Clean sigmoid on 15/15 architectures
LwF (output, soft)	0.73	32.6%	0.987	Clean sigmoid on 15/15 architectures
EWC (parameter, soft)	–	–	0.31	No reliable crossover on any architecture
SI (parameter, soft)	–	–	0.28	No reliable crossover on any architecture

FTR on this subset gives mean crossover 11.4, CV 20.5%, consistent with Section 6’s full-zoo result. LwF, fit with the same sigmoid procedure on all 15 architectures, gives a uniformly clean fit ($R^2 \geq 0.91$ on every architecture, median 0.987) with mean $\alpha^* = 0.73$, CV 32.6%, close to v1’s original $\alpha^* = 0.71$ on 5 CNN/ResNet architectures, now confirmed on $3\times$ as many architectures across 5 families, with correspondingly more (but still bounded) dispersion.

EWC and SI, by contrast, do not show a sigmoid transition on *any* of the 15 architectures when forced through the identical fitting procedure: median R^2 is 0.31 (EWC) and 0.28 (SI), with individual per-architecture point estimates spanning 4-5 orders of magnitude (h^* from 0.37 to 17,657 for EWC across architectures): not a location parameter with sampling noise, but a fit degenerating to noise because there is no transition to locate. This is a stronger confirmation of v1’s finding that EWC “shows no meaningful stability crossover” than v1’s own 5-architecture result: forcing the crossover-fitting machinery onto a flat curve produces wildly unstable estimates precisely *because* the underlying curve has no crossover, on $3\times$ the architecture diversity. This is consistent with Section 14’s structural argument that soft parameter-space penalties have no disengagement mechanism analogous to Eq. 8 and therefore no reason to produce a crossover at all.

9.1 A Pretrained-Backbone Sanity Check

Every result so far trains small models from scratch. This leaves open whether FTR’s advantage is an artifact of that regime or survives contact with how continual learning is actually practiced: fine-tuning a large pretrained backbone with a parameter-efficient adapter. We test this directly on ImageNet-pretrained ViT-B/16 (86M parameters) with LoRA (rank 8, $\alpha = 16$) applied to the MLP sublayers only, following Section C’s finding that `nn.MultiheadAttention`’s fused forward bypasses a wrapped `out_proj`; the base network is frozen, the classification head and LoRA adapters (0.84% of total parameters) are trainable. Split CIFAR-100, 10 tasks of 10 classes each, 200 samples/class, 2 epochs/task, 3 seeds. We compare FTR ($\varepsilon = 5.0$, $\eta_\lambda = 0.005$, $\lambda_{\text{init}} = 1.0$) against vanilla fine-tuning, EWC ($\lambda = 1000$), and LwF ($\alpha = 0.7$), each method’s hyperparameter fixed rather than swept, since a full crossover sweep on this scale is not the point of this check.

Table 6: Pretrained ViT-B/16 + LoRA, Split CIFAR-100, mean over 3 seeds. Forgetting and average accuracy are on the same $[0, 1]$ scale used throughout the paper.

Method	Forgetting	Avg. accuracy	Forgetting std (seeds)
Finetune (no constraint)	0.732	0.292	0.014
EWC ($\lambda = 1000$)	0.721	0.302	0.007
LwF ($\alpha = 0.7$)	0.611	0.374	0.018
FTR ($\varepsilon = 5.0$)	0.605	0.387	0.001

FTR gives both the lowest forgetting and the highest average accuracy of the four methods, ahead of LwF by a small margin and clearly ahead of EWC and unconstrained fine-tuning, which barely differ from each other. FTR’s forgetting is also the most stable across seeds (std 0.001, an order of magnitude tighter than LwF’s 0.018), notable given only 0.84% of the network’s parameters are trained. This is a single operating point, not a sweep, so it does not test architecture-invariance or the schedule confound; it tests only whether FTR’s advantage over softer constraints, established on small from-scratch models throughout this paper, transfers to a realistic parameter-efficient fine-tuning setting. It does.

10 KL-Direction Ablation

The constraint $D_{\text{KL}}(f_{\text{old}} \| f_{\theta}) \leq \varepsilon$ is asymmetric; we ablate the direction on CNN_W16 and CNN_W32 (dense ε grid, 3 seeds), comparing forward (paper default), reverse $D_{\text{KL}}(f_{\theta} \| f_{\text{old}})$, and symmetric Jensen-Shannon. All three directions fit cleanly ($R^2 = 0.83$ – 1.00). Forward and reverse are close: CNN_W16 gives $\varepsilon^* = 10.00$ (forward) vs. 10.32 (reverse); CNN_W32 gives 10.00 vs. 10.52, within 5% of each other. Jensen-Shannon is somewhat lower (9.49 and 7.79 respectively, a 5–22% decrease) but the same order of magnitude. KL direction is therefore a real but modest lever on ε^* : comparable to or smaller than momentum’s effect (Section 5) and far smaller than η_{λ} , λ_{init} , or steps/task.

11 Class-Incremental Variant

Every result so far uses the task-incremental protocol (Section 3): a head sized to classes-per-task, locally remapped labels, and an oracle task ID at evaluation. We additionally test the harder class-incremental protocol, a single shared 10-way head, no task ID, global labels throughout, on five width-family architectures (CNN_W8 through CNN_W48, 5 seeds, dense ε grid).

The phase transition survives (all five fits $R^2 \geq 0.91$), but at a substantially lower budget: class-incremental ε^* ranges 3.41–5.28, versus 10.62–11.72 for the same five architectures under the task-incremental protocol, a consistent 2–3 $\times$ reduction (ratio 0.32–0.50). This is expected and mechanistically sensible: with no oracle task ID, the model must keep every previously-seen class discriminable within one shared output space, a strictly harder requirement that a looser constraint cannot satisfy without additional forgetting. The qualitative finding, a sigmoid crossover exists and is comparable in relative dispersion across architectures (class-incremental CV 16.3% across these five architectures, versus 3.7% for the same five under task-incremental), holds under the harder protocol; the absolute budget required does not transfer between protocols and should not be assumed to.

12 Task-Structure Sensitivity: CIFAR-100 Granularity and Task Ordering

12.1 Classes-per-Task: a Third Instance of the Same Confound

Holding a fixed 20-class CIFAR-100 pool and varying classes/task $\in \{2, 4, 5, 10\}$ (10 architectures, 3 seeds, dense ε grid) tests whether ε^* depends on task granularity, directly probing the paper’s own claim that ε^* is set by task structure. The raw result is a large, mostly monotonic decrease in ε^* as classes/task increases: e.g. CNN_W16 gives $9.93 \rightarrow 4.43 \rightarrow 2.85 \rightarrow 2.14$ across $\text{cpt} = 2, 4, 5, 10$; 8 of 10 architectures show the same monotonic pattern.

This looks like a genuine task-structure effect, and was the intended result of the experiment. But Section 14.3’s own diagnostic obligates us to check for a schedule confound before reading it that way, and there is one: holding `max_per_class` and batch size fixed while classes/task increases also increases samples/task (and hence batches/task, and hence the constrained step budget N), purely mechanically. Using Eq. 9’s schedule term alone (ignoring any d_{unc} contribution): $\text{cpt} = 2$ gives $N \approx 21$, schedule term 9.52; $\text{cpt} = 4$ gives $N \approx 39$, term 5.13; $\text{cpt} = 5$ gives $N \approx 48$, term 4.17; $\text{cpt} = 10$ gives $N \approx 96$, term 2.08: the same monotonically decreasing pattern, comparable in magnitude to what we observe. We cannot, from this experiment as designed, separate a genuine task-structure effect from this third instance of the N -confound; a corrected version would need to hold samples/task (not samples/class) fixed as classes/task varies.

We ran that corrected version: 10 architectures, classes/task $\in \{2, 4, 5, 10\}$, samples/task now held fixed at 800 (`max_per_class` = 800/ cpt) so that N no longer varies mechanically with granularity, 3 seeds, same dense ε grid. The monotonic decrease disappears. Pooled across all 40 (architecture, granularity) cells, ε^* vs. classes/task gives $r = -0.174$ ($p = 0.283$), compared to the confounded design’s implied $r = -0.640$ ($p < 0.00001$) computed the same way: not a weaker version of the same effect, but no significant relationship at all. The $\text{cpt} = 2$ column is, however, its own outlier: curve fits there are visibly worse than at $\text{cpt} \geq 4$ ($R^2 = 0.32\text{--}0.89$ at $\text{cpt} = 2$ versus $R^2 \geq 0.89$ for every cell at $\text{cpt} \geq 4$), and two architectures (ResNet18_W8, ViT_Tiny) hit the same `curve_fit` bound-saturation artifact diagnosed in Section 6 at $\text{cpt} = 2$ specifically, plausibly because fixing samples/task at 800 while spreading it over only two classes leaves few examples per class and a noisier forgetting curve. Restricting to the reliably-fit range $\text{cpt} \in \{4, 5, 10\}$ (every cell $R^2 \geq 0.89$), the pooled correlation flips sign but remains non-significant: $r = +0.265$ ($p = 0.157$, $n = 30$), and the mean within-architecture Spearman correlation across this range is $+0.45$, weakly positive rather than negative. We do not read this as evidence of a genuine increasing effect either, given $p = 0.157$ and $n = 10$ architectures; the honest summary is that once the schedule confound is removed, classes-per-task shows no reliable, significant relationship with ε^* over the range we can fit reliably (4–10 classes/task), and the striking monotonic decrease in the original design was overwhelmingly the N -confound, not a genuine task-structure effect: the same lesson as Sections 5 and 14.3, discovered a third time and this time directly confirmed rather than left inconclusive.

12.2 Task Ordering: a Genuine, Non-Schedule Source of Variability

On 6 architectures, 3 random class-to-task permutations (identical N across permutations, so this test is *not* subject to the confound above), ε^* varies more than seed noise alone would predict: CNN_W32, for the two reliably-fit orderings ($R^2 \geq 0.97$), gives $\varepsilon^* = 36.4$ and 8.6 , a $4.2\times$ range; CNN_W16 gives 22.3 and 6.9 ($3.2\times$); ResNet18_W8, fit reliably on all three orderings, gives 9.8 , 12.6 , 8.2 (a tighter $1.5\times$ range). One ordering (seed 1001) fits poorly on several architectures (R^2 as low as 0.31), so these point estimates should be read as indicating a real, substantial effect of which

specific classes are grouped into which task, not as precisely located crossovers. Unlike the diagnostic sweep (Section 5), this variability survives with the schedule held completely fixed, and is a genuine task-structure effect: which classes end up adjacent in the task sequence measurably changes d_{unc} , the natural drift rate a new task induces. v1 tested exactly one task ordering throughout; this result indicates that choice was not neutral.

13 S-Invariance: Collapsing the Schedule Confound to a Single Ratio

Section 5 shows η_λ , λ_{init} , and steps/task each independently move ε^* by more than an order of magnitude. Taken at face value, this leaves a practitioner with a three-dimensional sensitivity problem: any reported ε^* is conditional on three separately-chosen numbers, with no obvious way to report or budget for that dependence short of a full grid search. Equation 9’s schedule term, $\lambda_{\text{init}}/(\eta_\lambda N)$, suggests a way out: if this specific ratio, which we call S , is what actually drives the schedule’s effect on ε^* , then the three-dimensional problem collapses to one dimension, and holding S fixed while varying η_λ , λ_{init} , and N individually should leave ε^* approximately unchanged even though each of the three still moves it by an order of magnitude on its own.

13.1 Design

On five architectures spanning four of the five families (two CNN widths, ResNet18_W16, MLP_H128, ViT_Tiny; MLP-Mixer was not included in this sub-zoo), we ran six conditions, all epoch/task schedule-matched at 4 epochs/task: the Section 5 baseline ($\eta_\lambda = 0.005$, $\lambda_{\text{init}} = 1.0$, $N = 24$, giving $S = 8.33$); two conditions that scale η_λ and λ_{init} together by the same factor so S is unchanged ($\times 5$ and $\times 0.2$, still $S = 8.33$); two conditions that instead reach the same $S = 8.33$ by doubling the step budget N (epochs/task forced to 7) while adjusting η_λ or λ_{init} to compensate; and one negative control that changes only η_λ by $\times 5$ with nothing else adjusted, giving $S = 1.67$, a genuine $5\times$ change in S itself. If S is the right invariant, the first five conditions (identical S , three different routes to it) should cluster, and the sixth (different S) should not. Three seeds per (architecture, condition, ε) cell, dense ε grid. Four of the resulting 30 cells initially hit the same curve_fit bound-saturation issue diagnosed in Section 6 (three baselines and one scale-up condition, each pinned at $50e = 54.37$ on this experiment’s narrower grid); we reran exactly those four on a wider grid extending to $\varepsilon = 200$ (5 seeds, 1,000 bootstrap resamples) to locate their true crossovers rather than report degenerate fits.

Table 7: ε^* under five different routes to the same schedule ratio $S = \lambda_{\text{init}}/(\eta_\lambda N) = 8.33$, plus one negative control with $S = 1.67$ (a genuine $5\times$ change in S), architecture and task held fixed. CV is computed over the five $S = 8.33$ columns only. [†]Relocated on the wider grid after the narrow-grid fit saturated; see text.

Architecture	baseline $S = 8.33$	scale $\times 5$ $S = 8.33$	scale $\times 0.2$ $S = 8.33$	steps-a $S = 8.33$	steps-b $S = 8.33$	CV (5 cols)	mismatched $S = 1.67$
CNN_W16	10.08	12.42	14.35	8.72	10.16	17.9%	2.73
CNN_W32	15.79	13.61	5.88	13.54	10.10	29.4%	1.06
ResNet18_W16	11.26 [†]	12.91	22.02	10.14	8.96	35.7%	1.45
MLP_H128	12.84 [†]	20.05	8.92	8.14	9.69	36.6%	3.02
ViT_Tiny	9.99 [†]	18.21 [†]	18.99	8.29	7.86	38.7%	1.18




Figure 8: For each architecture, the five $S = 8.33$ conditions (circles) cluster together while the $S = 1.67$ mismatched control (X) sits well below the cluster on every architecture, log scale. The residual spread within each cluster (vertical bar marks the cluster mean) is real but small relative to the gap to the mismatched control.

13.2 Results

Table 7 and Figure 8 show a clear qualitative split. Within the five $S = 8.33$ columns, ϵ^* ranges roughly 6–22 per architecture (CV 17.9%–38.7%, mean 31.7% across the five architectures): a real residual spread, not exact invariance, but far tighter than what any single schedule parameter produces on its own. The mismatched control, which changes only η_λ by $5\times$ (an S change of the same $5\times$ magnitude as the within-group perturbations, but applied to S itself rather than cancelled against it), collapses ϵ^* to 9–25% of the matched-group mean (a 4–11 $\times$ reduction, consistent in direction and magnitude with the $\eta_\lambda \times 5$ row of Table 1). S is not a perfect invariant of ϵ^* , in the way architecture is approximately once the schedule itself is held fixed (Section 7), but it captures the dominant share of the schedule’s effect: three independently-chosen hyperparameters that each move ϵ^* by more than an order of magnitude on their own (Section 5) reduce, once combined into S , to a quantity whose value alone predicts ϵ^* to within roughly a factor of 2–3 across five different architectures and three different routes to the same S .

13.3 Practical Upshot

This converts an intractable three-dimensional sensitivity into a single reportable ratio. A practitioner adopting FTR, or any Lagrangian-dual constrained continual-learning method with a periodically reset dual variable, does not need to separately justify η_λ , λ_{init} , and steps/task: reporting $S = \lambda_{\text{init}}/(\eta_\lambda N)$ alongside ϵ^* captures most of what those three numbers jointly determine, and two runs with matched S but different individual hyperparameters should be expected to land within roughly a factor of 2–3 of each other, not the order-of-magnitude range Section 5 shows is possible when S itself is left to vary. We do not claim this closes the schedule-dependence problem entirely, the residual 18–39% within- S spread is real and architecture-dependent, but it is a substantial, prac-

tically useful reduction in dimensionality, and, to our knowledge, the first time this specific ratio has been identified and tested as the operative quantity in Lagrangian-dual constrained continual learning.

14 Interpretation and Theory

14.1 A Back-of-Envelope Account of ε^*

Section 5 showed ε^* is dominated by the optimizer’s own schedule. We now give an approximate, testable account of *why*, and of why architecture-independence nonetheless survives at a fixed schedule.

Consider a task transition with warmup complete, N constrained gradient steps remaining in the task’s budget, and dual-ascent momentum ignored for this first-order treatment ($\rho = 0$; the EMA in Eq. 4 otherwise smooths but does not change the sign of the drift). If the model’s drift under the *current* λ runs at an approximately constant rate $\bar{d}(\lambda)$ per step, the violation at each step is $\bar{d}(\lambda_t) - \varepsilon$, and

$$\lambda_{t+1} \approx \lambda_t + \eta_\lambda (\bar{d}(\lambda_t) - \varepsilon). \quad (7)$$

Write $d_{\text{unc}} = \bar{d}(0)$, the drift rate the model would produce under *no* constraint (i.e. as $\lambda \rightarrow 0$), a property of the task and architecture, not the optimizer. If $\varepsilon > d_{\text{unc}}$, the violation is negative even at $\lambda = 0$, so λ decays monotonically from λ_{init} and reaches 0 after approximately

$$t^*(\varepsilon) \approx \frac{\lambda_{\text{init}}}{\eta_\lambda (\varepsilon - d_{\text{unc}})} \quad (8)$$

steps (treating $\bar{d}(\lambda_t) \approx d_{\text{unc}}$ near $\lambda \approx 0$, where the model is close to unconstrained). The task-level crossover occurs where the constraint’s disengagement time equals the step budget, $t^*(\varepsilon^*) = N$, giving

$$\varepsilon^* \approx d_{\text{unc}} + \frac{\lambda_{\text{init}}}{\eta_\lambda N}. \quad (9)$$

This single approximate expression explains both findings at once. First, the second term is a pure function of the dual-ascent schedule (λ_{init} , η_λ , N) and does not depend on architecture at all, consistent with Section 5’s finding that these three quantities dominate ε^* and move it by more than an order of magnitude. Second, because this schedule term is *identical across every architecture trained with the same hyperparameters*, it automatically produces architecture-independence in ε^* whenever the schedule is held fixed, regardless of how much d_{unc} (the genuinely architecture-dependent term) varies, explaining why Section 6’s expanded zoo still shows family-level clustering rather than architecture-independence being destroyed outright: the clustering is exactly the residual d_{unc} signal riding on top of a shared, schedule-determined baseline.

Equation 9 is only approximate (it ignores momentum, the EMA smoothing constant β , and treats $\bar{d}$ as constant in λ rather than solving the coupled primal-dual dynamics self-consistently), so we treat it as a mechanistic account, not a closed-form prediction, and check it against data rather than asserting it. With $\lambda_{\text{init}} = 1.0$, $\eta_\lambda = 0.005$: at $N = 24$ (the diagnostic’s baseline, 4 epochs/task minus 1 warmup epoch, 8 batches/epoch), the schedule term is $1/(0.005 \times 24) = 8.33$; at $N = 56$ (the $\times 2$ epochs/task condition), it is $1/(0.005 \times 56) = 3.57$, versus the observed $\varepsilon^* = 3.25\text{--}3.66$ (Table 1), within 10% using only the schedule term, no architecture correction at all. At $N = 24$ (baseline), the schedule term 8.33 under-predicts the observed 12.45–13.09 by a d_{unc} correction of 4.1–4.8, comparable in magnitude to the schedule term itself rather than negligible, so while the schedule term dominates *sensitivity* (Section 5’s core finding), it does not always dominate the

absolute value, and the architecture-dependent correction is real and worth taking seriously, which brings us to a confound in our own zoo design.

14.2 Convergence-Theoretic Grounding for the Schedule Term

The $1/(\eta_\lambda N)$ scaling in Eq. 9 is not an arbitrary fit; it is the signature of a standard fact about dual ascent. For a concave dual function optimized by projected subgradient ascent with fixed step size η_λ , the sub-optimality of the dual iterate after N steps decays at rate $O(1/(\eta_\lambda N))$ under standard step-size and boundedness conditions [Boyd and Vandenberghe, 2004, Nedić and Ozdaglar, 2009]; this is the same $O(1/T)$ behavior that underlies our Eq. 8’s disengagement-time argument, just stated as a convergence-rate bound rather than a hitting-time calculation. Chamon et al. [2023] extend exactly this style of guarantee, near-PAC learnability and a bounded empirical duality gap under dual ascent, to the non-convex losses that neural-network training loss functions actually are, which is the setting FTR’s task loss and constraint fall into. Read together, their result establishes that a dual-ascent-enforced constraint on a non-convex loss converges to near-feasibility at a rate governed by η_λ and the step budget rather than by problem-specific constants alone; our diagnostic sweep (Section 5) is, in effect, an empirical demonstration that this generic rate dependence is large enough, in the specific regime FTR operates in, to dominate the architecture-dependent constants it is usually assumed to be secondary to. We do not derive Eq. 9 as a formal corollary of their bound (that would require tracking the exact non-convex duality-gap constants through to a task-transition hitting time, which we leave as a natural extension), but the qualitative scaling, and the fact that Section 13 finds the schedule collapses to a single ratio rather than three independent effects, is exactly what generic dual-ascent convergence theory predicts and is not, to our knowledge, a connection the Lagrangian-dual continual-learning literature has previously drawn out explicitly.

14.3 A Self-Inflicted Version of the Same Confound

Equation 9 makes a specific, checkable prediction about our own experimental design. The ViT and MLP-Mixer families in Section 4 were assigned 5 epochs/task rather than the 4 used for every other family (a well-intentioned accommodation for harder optimization at the time the zoo was built), giving them $N = 32$ rather than $N = 24$ constrained steps. Equation 9 predicts this alone should shift their ε^* down by approximately $1/(0.005 \times 24) - 1/(0.005 \times 32) = 8.33 - 6.25 = 2.08$ relative to an epoch-matched comparison, *before any genuine architectural difference is considered*.

We tested this directly: re-running all four affected architectures (ViT_Tiny, Mixer_Tiny, and the same mismatch on their wider siblings ViT_Small and Mixer_Small) with epochs forced to 4 (matching the CNN/ResNet/MLP family, 5 seeds, same dense ε grid as Section 6). All four shift up substantially (Table 8): ViT_Tiny moves from $\varepsilon^* = 7.31$ (5 epochs/task) to 10.18 [9.78, 10.88] (4 epochs/task, $R^2 = 0.93$), a shift of +2.86, close to the predicted +2.08. Mixer_Tiny moves from 6.10 to 10.41 [8.76, 11.75] ($R^2 = 0.99$), a shift of +4.30. ViT_Small moves from 7.40 to 12.43 [11.25, 15.58] ($R^2 = 0.98$), a shift of +5.03. Mixer_Small moves from 6.11 to 10.63 [9.89, 11.47] ($R^2 = 0.98$), a shift of +4.52. All four shifts are in the predicted direction and, for three of four, exceed the back-of-envelope prediction of +2.08, consistent with the architectures also differing in genuine optimization difficulty at matched epoch count, on top of the schedule effect.

Epoch-matched, all four land at ≈ 10.2 – 12.4 , inside the CNN/ResNet/MLP family’s range (8.35–15.98 across its 26 architectures) rather than systematically below it. The family clustering reported in Section 7.1 was therefore substantially, possibly almost entirely, an artifact of our own epoch/task mismatch, not a genuine effect of representational paradigm (convolution vs. attention vs. token-mixing). We report this plainly rather than let the earlier, uncorrected framing stand,

Table 8: Epoch-matched control: re-running the four 5-epoch/task architectures at 4 epochs/task, matching every other family.

Architecture	ε^* (5 ep/task, as recorded)	ε^* (4 ep/task, matched)	shift
ViT_Tiny	7.31	10.18 [9.78, 10.88]	+2.86
Mixer_Tiny	6.10	10.41 [8.76, 11.75]	+4.30
ViT_Small	7.40	12.43 [11.25, 15.58]	+5.03
Mixer_Small	6.11	10.63 [9.89, 11.47]	+4.52

since it is exactly the kind of check Section 5 argues the field should be running as standard practice, and we caught it only by holding ourselves to that standard against our own zoo. Sections 7.1 and 7.2 are revised accordingly, substituting these four epoch-matched values for the confounded ones throughout the rest of the paper.

14.4 Why Curvature Still Fails to Predict ε^*

Equation 9 also clarifies why curvature-based normalization fails (Section 7): ε^* 's dominant term is a property of the optimizer, and its secondary term (d_{unc}) is the drift rate of *unconstrained* gradient descent measured in the KL-divergence output-space metric, not a parameter-space curvature quantity. A model's Hessian trace or Fisher trace describes local loss-landscape curvature in parameter space; d_{unc} describes how fast the model's *output distribution* moves under gradient descent on a new task, which depends on parameter-space curvature only through a Jacobian that varies with architecture in a way that (per Section 7.4's adequately-powered $n = 30$ test) none of Hessian trace, Fisher trace, spectral norm, d_{eff} , or parameter count predicts. This is consistent with, and does not contradict, Lanzillotta et al. [2023]'s formal result that sharper minima require tighter *parameter*-space constraints: that result concerns parameter-space penalties (EWC-like), which is a different constraint geometry from the function-space KL ball studied here.

14.5 Why Parameter-Space Methods Fail to Show Transitions

EWC minimizes the unconstrained objective $\mathcal{L}_{\text{task}}(\theta) + \frac{\lambda}{2}(\theta - \theta_{\text{old}})^\top F(\theta - \theta_{\text{old}})$, a soft penalty with no hard constraint boundary and no dual variable that can be driven to zero mid-task the way λ can here. As λ increases, forgetting decreases smoothly; there is no analogue of Eq. 8's disengagement time, and hence no mechanism to produce a crossover. This is a structural property of soft penalties versus hard constraints with adaptive duals, independent of the diagnostic finding above.

15 Limitations

This work retains several limitations from v1 and adds new ones surfaced by this revision itself.

1. **The numeric value of ε^* is schedule-conditional, not a hyperparameter-free constant.** This is the paper's central revision, not merely a caveat: Section 5 shows ε^* moves by more than an order of magnitude under simple, reasonable changes to η_λ , λ_{init} , or steps/task. Any reuse of a reported ε^* value from this paper (or any Lagrangian-dual constrained continual-learning method) on a different codebase should first check sensitivity to these three quantities, not assume architecture is the only relevant axis.

2. **Single dataset.** All main experiments use CIFAR-10; the CIFAR-100 granularity sweep (Section 12) holds the dataset fixed and varies only task construction. Tiny-ImageNet, CORe50, and CLEAR-style benchmarks remain untested.
3. **Modest scale.** Moving to GPU compute (from v1’s CPU-only setup) enabled a $3.75\times$ larger architecture zoo and more seeds, but model sizes ($\leq 2.8\text{M}$ parameters) and per-task sample counts (1,000/class) remain small by contemporary standards.
4. **The theoretical account (Section 14) is approximate, not a closed-form derivation.** Equation 9 ignores dual-ascent momentum and treats the drift rate as constant in λ ; it matches the epochs/task sensitivity closely but under-predicts the baseline value by an amount comparable to the schedule term itself, so it should be read as a mechanistic account that generates testable predictions (Section 14.3), not an exact formula.
5. **Aggregate and per-task-profile metrics only go so far.** We report new-task accuracy and per-task forgetting profiles (Section 7.5), an improvement on v1’s aggregate-only reporting, but do not analyze per-class or per-sample forgetting patterns.
6. **No accuracy/SOTA comparison.** Unchanged from v1: this work is diagnostic and mechanistic, not a proposal for a state-of-the-art continual-learning method, and FTR without replay does not exceed LwF on average accuracy.
7. **Curvature measurement required a backend workaround for one family.** The Hessian trace / spectral norm estimators require a twice-differentiable loss surface; PyTorch’s fused scaled-dot-product-attention kernel does not support this by default, which initially excluded the ViT family from curvature measurement entirely. Forcing the math-mode SDPA backend for curvature measurement only (Appendix B) resolved this, and the correlation results in Section 7 use the full $n = 30$ zoo.
8. **One architecture’s crossover remains imprecisely determined even after correction.** ResNetLite_W8_NoBN’s wide-grid-relocated bootstrap CI is [6.87, 97.02] (Section 6); extending the search grid resolved the exact bound-pinning artifact but not the underlying measurement noise, which a purely ε -grid-based fix cannot address. More seeds, rather than a wider grid, are the likelier next step for this specific architecture.
9. **The classes-per-task=2 condition is unreliable in the corrected CIFAR-100 granularity experiment.** Holding samples/task fixed at this granularity leaves few examples per class; fit quality there is visibly worse than at classes/task ≥ 4 (Section 12.1), and we do not draw conclusions from that specific column.

16 Conclusion

We set out to test a specific, previously unexamined confound in Lagrangian-dual constrained continual learning: because the dual variable λ resets at every task boundary, the observed stability crossover of a function-space KL constraint could in principle be an artifact of the constrained optimizer’s own schedule rather than a property of the task or architecture it is applied to. It is. Holding architecture and task fixed, the dual learning rate, the initial multiplier, and the per-task step budget each move ε^* by more than an order of magnitude, more than a 75-fold range in architecture size moves it, more than any other axis we tested. A back-of-envelope mechanistic account (Eq. 9) explains both this sensitivity and why architecture-independence survives it: the

dominant term in ε^* is a pure function of the schedule, shared identically across every architecture trained with the same hyperparameters, while the smaller, genuinely architecture-dependent term is what curvature-based normalization has always tried and failed to capture.

That failure is real and, at the sample size this revision affords, now properly powered: across 30 architectures spanning five representational families, no tested curvature statistic predicts ε^* , extending v1’s claim from an $n = 8$ test too underpowered to have supported it into an adequately powered one that does. What curvature does not predict, our own zoo design initially seemed to: ViT and MLP-Mixer sat systematically lower than CNN/ResNet/MLP, until applying the paper’s own central diagnostic to the paper’s own architecture zoo showed most of that gap was a second, self-inflicted instance of the identical schedule confound. We report this correction plainly, because catching it is a demonstration of the paper’s thesis, not a footnote to it: this class of confound is easy to introduce by accident and easy to miss without deliberately checking for it, and it recurred a third time in an experiment explicitly designed to probe task-structure sensitivity (Section 12.1).

The practical guidance this leaves practitioners of Lagrangian-relaxed, dual-ascent continual learning methods is the inverse of v1’s: architecture-specific tuning of the stability budget is not necessary, but the dual-ascent schedule is not a free implementation detail either, and checking sensitivity to η_λ , λ_{init} , and steps-per-task belongs in the standard reporting protocol for any method in this family, ours included. Task ordering and the choice between task-incremental and class-incremental evaluation are further, genuine sources of variation this paper surfaces that v1’s single fixed ordering and single protocol could not. Open questions remain: what remains of the schedule-corrected residual family effect (Section 7.2 finds it small but not exactly zero, $\text{ICC} \approx 0.05$) traces to; whether d_{unc} in Eq. 9 can be derived rather than measured; and whether the same confound structure applies to other Lagrangian-dual methods in continual learning beyond FTR [Elenter et al., 2023, Chamon et al., 2023]. We see the diagnostic protocol demonstrated here, vary the constrained optimizer’s own hyperparameters before attributing an observed threshold to the model or the task, as the paper’s most transferable contribution, independent of where any particular numeric value of ε^* ultimately lands.

References

- Shun-ichi Amari. Natural gradient works efficiently in learning. *Neural Computation*, 10(2):251–276, 1998.
- Ang Bian, Wei Li, Hangjie Yuan, Chengrong Yu, Mang Wang, Zixiang Zhao, Aojun Lu, Pengliang Ji, and Tao Feng. Make continual learning stronger via C-Flat. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2404.00986.
- Stephen Boyd and Lieven Vandenbergh. *Convex Optimization*. Cambridge University Press, 2004.
- Pietro Buzzega, Matteo Boschini, Angelo Porrello, and Simone Calderara. Dark experience for general continual learning: a strong, simple baseline. *Advances in Neural Information Processing Systems (NeurIPS)*, 33, 2020.
- Luiz F. O. Chamon, Santiago Paternain, Miguel Calvo-Fullana, and Alejandro Ribeiro. Constrained learning with non-convex losses. *IEEE Transactions on Information Theory*, 69(3):1739–1760, 2023. arXiv:2103.05134.
- Danruo Deng, Guangyong Chen, Jianye Hao, Qiong Wang, and Pheng-Ann Heng. Flattening sharpness for dynamic gradient projection memory benefits continual learning. *Advances in Neural Information Processing Systems (NeurIPS)*, 34, 2021.

- Felix Draxler, Kambis Veschgini, Manfred Salmhofer, and Fred Hamprecht. Essentially no barriers in neural network energy landscape. In *International Conference on Machine Learning (ICML)*, pages 1309–1318, 2018.
- Juan Elenter, Navid NaderiAlizadeh, Tara Javidi, and Alejandro Ribeiro. Primal dual continual learning: Balancing stability and plasticity through adaptive memory allocation. *arXiv preprint arXiv:2310.00154*, 2023.
- Juan Elenter, Luiz F. O. Chamon, and Alejandro Ribeiro. Near-optimal solutions of constrained learning problems. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2403.11844.
- Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. In *International Conference on Learning Representations (ICLR)*, 2019.
- Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis. In *International Conference on Machine Learning (ICML)*, pages 3259–3269, 2020.
- Timur Garipov, Pavel Izmailov, Dmitrii Podoprikin, Dmitry P. Vetrov, and Andrew G. Wilson. Loss surfaces, mode connectivity, and fast ensembling of DNNs. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 31, 2018.
- Behrooz Ghorbani, Shankar Krishnan, and Ying Xiao. An investigation into neural net optimization via Hessian eigenvalue density. In *International Conference on Machine Learning (ICML)*, pages 2232–2241, 2019.
- Elad Hazan. *Introduction to Online Convex Optimization*. MIT Press, 2nd edition, 2019.
- Ignacio Hounie, Alejandro Ribeiro, and Luiz F. O. Chamon. Resilient constrained learning. *Advances in Neural Information Processing Systems (NeurIPS)*, 36, 2023. arXiv:2306.02426.
- Nitish Shirish Keskar, Dheevatsa Mudigere, Jorge Nocedal, Mikhail Smelyanskiy, and Ping Tak Peter Tang. On large-batch training for deep learning: Generalization gap and sharp minima. In *International Conference on Learning Representations (ICLR)*, 2017.
- Snigdha Chandan Khilar. Continual learning as a multiphase moving-boundary problem. *arXiv preprint arXiv:2606.01863*, 2026. unreviewed preprint.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*, 2015.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- Giulia Lanzillotta, Sidak Pal Singh, Benjamin F. Grewe, and Thomas Hofmann. Towards guarantees for parameter isolation in continual learning. *arXiv preprint arXiv:2310.01165*, 2023.

- Alex Lewandowski, Hidenori Tanaka, Dale Schuurmans, and Marlos C. Machado. Directions of curvature as an explanation for loss of plasticity. *arXiv preprint arXiv:2312.00246*, 2023.
- Alex Lewandowski, Michał Bortkiewicz, Saurabh Kumar, András György, Mateusz Ostaszewski, Dale Schuurmans, and Marlos C. Machado. Learning continually by spectral regularization. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2406.06811.
- Hao Li, Zheng Xu, Gavin Taylor, Christoph Studer, and Tom Goldstein. Visualizing the loss landscape of neural nets. *Advances in Neural Information Processing Systems (NeurIPS)*, 31, 2018.
- Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(12):2935–2947, 2017.
- David Lopez-Paz and Marc’Aurelio Ranzato. Gradient episodic memory for continual learning. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 2017.
- James Martens and Roger Grosse. Optimizing neural networks with Kronecker-factored approximate curvature. In *International Conference on Machine Learning (ICML)*, pages 2408–2417, 2015.
- Michael McCloskey and Neal J. Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 24:109–165, 1989.
- Angelia Nedić and Asuman Ozdaglar. Subgradient methods for saddle-point problems. *Journal of Optimization Theory and Applications*, 142(1):205–228, 2009.
- Roman Novak, Yasaman Bahri, Daniel A. Abolafia, Jeffrey Pennington, and Jascha Sohl-Dickstein. Sensitivity and generalization in neural networks: an empirical study. In *International Conference on Learning Representations (ICLR)*, 2018.
- Pingbo Pan, Siddharth Swaroop, Alexander Immer, Runa Eschenhagen, Richard E. Turner, and Mohammad Emtiyaz Khan. Continual deep learning by functional regularisation of memorable past. *Advances in Neural Information Processing Systems (NeurIPS)*, 33, 2020. arXiv:2004.14070.
- Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*, 2022.
- Alexander Rakhlin, Karthik Sridharan, and Ambuj Tewari. Online learning: Random averages, combinatorial parameters, and learnability. *Journal of Machine Learning Research*, 13(1):2599–2633, 2012.
- Anthony Robins. Catastrophic forgetting, rehearsal and pseudorehearsal. *Connection Science*, 7(2): 123–146, 1995.
- John Schulman, Sergey Levine, Pieter Abbeel, Michael Jordan, and Philipp Moritz. Trust region policy optimization. In *International Conference on Machine Learning (ICML)*, pages 1889–1897, 2015.
- Shai Shalev-Shwartz. Online learning and online convex optimization. *Foundations and Trends in Machine Learning*, 4(2):107–194, 2012.
- Hanul Shin, Jung Kwon Lee, Jaehong Kim, and Jiwon Kim. Continual learning with deep generative replay. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 2017.

Keiichi Tamai, Tsuyoshi Okubo, Truong Vinh Truong Duy, Naotake Natori, and Synge Todo. Universal scaling laws of absorbing phase transitions in artificial deep neural networks. *Physical Review Research*, 7:033072, 2025. arXiv:2307.02284.

Michalis K. Titsias, Jonathan Schwarz, Alexander G. de G. Matthews, Razvan Pascanu, and Yee Whye Teh. Functional regularisation for continual learning with Gaussian processes. In *International Conference on Learning Representations (ICLR)*, 2020. arXiv:1901.11356.

Ping Wang. Grokking as dimensional phase transition in neural networks. *arXiv preprint arXiv:2604.04655*, 2026. unreviewed preprint.

Friedemann Zenke, Ben Poole, and Surya Ganguli. Continual learning through synaptic intelligence. In *International Conference on Machine Learning (ICML)*, pages 3987–3995, 2017.

A Full Architecture Zoo and Crossover Table

Table 9 reports ε^* , its 95% bootstrap CI, sigmoid R^2 , and sharpness k for all 30 architectures, sorted by parameter count, exactly as originally recorded (5 epochs/task for ViT/Mixer, narrow ε grid for every row) for a complete, unedited record of the raw campaign. Six of these thirty rows are superseded elsewhere in the paper and should not be used for any downstream analysis: ViT_Tiny, Mixer_Tiny, ViT_Small, and Mixer_Small by their epoch-matched values (Table 8), and ResNet18_W16 and ResNetLite_W8_NoBN by their wide-grid-relocated values (Section 6). Table 4, the hierarchical model of Section 7.2, and every downstream analysis use the corrected values for these six architectures. Table 10 reports the corresponding Hessian trace, Fisher trace, spectral norm, and d_{eff} used in Section 7’s correlation analysis.

Table 9: Full 30-architecture crossover table (CIFAR-10, dense ε sweep, 5 seeds per architecture, 10 for CNN_W16/CNN_W32).

Architecture	Family	Params	ε^*	95% CI	R^2	k
CNN_W8_NoBN	cnn	36,866	21.45	[18.57, 30.88]	0.986	2.35
CNN_W8	cnn	36,946	11.72	[10.11, 14.64]	0.996	2.38
Mixer_Tiny	mixer	45,442	6.10	[5.88, 6.50]	0.972	4.04
ViT_Tiny	vit	74,562	7.31	[6.35, 8.03]	0.974	4.19
ResNetLite_W8_NoBN	resnet	77,122	15.98	[11.80, 135.91]	0.783	2.67
ResNetLite_W8	resnet	77,482	10.46	[8.18, 13.41]	0.962	3.75
CNN_W16	cnn	80,418	10.80	[10.03, 11.48]	0.990	4.04
CNN_D4_W32	cnn	126,850	11.62	[10.13, 13.32]	0.970	2.98
CNN_W24	cnn	130,802	10.65	[9.81, 11.69]	0.977	4.14
CNN_D5_W32	cnn	135,042	10.10	[9.24, 11.35]	0.978	3.42
Mixer_Small	mixer	136,466	6.11	[5.63, 6.53]	0.943	4.08
ResNet18_W8	resnet	175,882	13.76	[10.98, 26.73]	0.901	2.16
CNN_W32_NoBN	cnn	187,778	11.70	[10.85, 12.36]	0.987	3.71
CNN_W32	cnn	188,098	10.62	[9.33, 11.90]	0.988	3.38
CNN_D3_W32	cnn	188,098	11.36	[10.83, 12.34]	0.992	4.39
ResNetLite_W16	resnet	307,794	11.41	[8.93, 14.77]	0.968	2.26
ViT_Small	vit	310,562	7.40	[6.65, 8.38]	0.980	2.17
CNN_W48	cnn	323,426	11.08	[9.96, 12.45]	0.984	4.94
MLP_H128	mlp	410,114	14.56	[13.74, 15.37]	0.994	2.10
CNN_W64_NoBN	cnn	485,762	11.64	[10.99, 12.36]	0.967	4.57
CNN_W64	cnn	486,402	11.18	[10.37, 12.17]	0.985	3.31
CNN_D2_W32	cnn	544,258	10.76	[9.98, 11.65]	0.983	3.57
ResNetLite_W24	resnet	690,938	8.35	[6.60, 10.64]	0.944	3.01
ResNet18_W16	resnet	700,434	14.59	[9.65, 135.91]	0.937	1.63
MLP_H256	mlp	852,994	12.52	[11.82, 13.35]	0.991	2.70
CNN_W96	cnn	895,298	11.10	[10.56, 11.88]	0.974	4.23
ResNetLite_W32	resnet	1,226,914	8.39	[6.56, 9.71]	0.977	2.66
CNN_W128	cnn	1,414,786	10.43	[9.85, 10.94]	0.979	3.84
ResNet18_W24	resnet	1,573,658	12.15	[9.73, 25.01]	0.927	2.29
ResNet18_W32	resnet	2,795,554	8.88	[7.45, 11.84]	0.947	2.33

Table 10: Curvature statistics for all 30 architectures (mean over 5 seeds, measured after Task 1 training, matching the protocol of Section 4.1).

Architecture	Params	$\text{tr}(H)$	$\text{tr}(F)$	Spectral norm	d_{eff}
CNN_W8_NoBN	36,866	118.9	0.926	78.8	1.51
CNN_W8	36,946	227.6	1.126	55.5	4.98
Mixer_Tiny	45,442	124.2	0.253	18.5	7.05
ViT_Tiny	74,562	312.4	3.776	132.8	2.30
ResNetLite_W8_NoBN	77,122	93.0	1.069	70.1	1.36
ResNetLite_W8	77,482	951.8	2.238	133.1	8.74
CNN_W16	80,418	318.6	2.312	62.4	5.46
CNN_D4_W32	126,850	429.2	4.340	65.3	7.25
CNN_W24	130,802	309.6	2.098	62.6	5.33
CNN_D5_W32	135,042	391.7	1.577	39.5	10.77
Mixer_Small	136,466	145.1	0.394	25.9	5.51
ResNet18_W8	175,882	3110.5	2.753	235.5	14.65
CNN_W32_NoBN	187,778	79.7	1.361	50.4	1.62
CNN_W32	188,098	305.7	1.800	48.6	6.50
CNN_D3_W32	188,098	307.8	1.825	48.9	6.49
ResNetLite_W16	307,794	992.6	2.774	106.1	9.99
ViT_Small	310,562	110.0	9.427	69.7	1.49
CNN_W48	323,426	203.5	1.085	35.8	5.70
MLP_H128	410,114	129.4	1.283	27.9	4.70
CNN_W64_NoBN	485,762	98.0	5.718	52.8	1.89
CNN_W64	486,402	180.3	1.064	32.4	5.75
CNN_D2_W32	544,258	229.1	2.703	56.9	3.91
ResNetLite_W24	690,938	794.0	1.884	74.6	13.91
ResNet18_W16	700,434	1235.8	3.118	127.8	13.81
MLP_H256	852,994	161.8	1.153	37.7	4.40
CNN_W96	895,298	114.7	1.192	34.8	3.27
ResNetLite_W32	1,226,914	556.1	1.368	48.6	12.93
CNN_W128	1,414,786	147.5	0.659	40.3	3.72
ResNet18_W24	1,573,658	1013.2	2.920	55.0	19.87
ResNet18_W32	2,795,554	483.2	3.169	67.7	8.06

B SDPA Backend Workaround for Curvature Measurement

The Hutchinson Hessian-trace estimator and the power-iteration spectral-norm estimator both require a second derivative of the loss with respect to parameters (`torch.autograd.grad` with `create_graph=True` called twice). PyTorch’s default scaled-dot-product-attention kernel (dispatched internally by `nn.MultiheadAttention`, used by the ViT family) selects a fused backend whose backward pass is not itself twice-differentiable, raising `derivative for aten::_scaled_dot_product_efficient_attention_backward is not implemented` and failing 100% of curvature measurements for ViT_Tiny and ViT_Small in our initial run. Wrapping only the curvature-measurement forward/backward passes in `torch.nn.attention.sdpa_kernel([SDPBackend.MATH])` (or the equivalent `torch.backends.cuda.sdp_kernel` context manager on older PyTorch versions) forces the plain, twice-differentiable attention implementation for those calls only; ordinary training is unaffected and continues to use the faster fused kernel. This recovered all 30 architectures for the correlation analysis in Section 7.

C LoRA Placement for the Pretrained-Backbone Check

Wrapping `nn.MultiheadAttention`’s `out_proj` submodule with a LoRA adapter has no effect at inference or training time: PyTorch’s `nn.MultiheadAttention.forward` does not call `out_proj.forward()` internally, it reads `out_proj.weight` and `out_proj.bias` directly inside a single fused functional attention call, so any wrapper module placed on `out_proj` is silently bypassed and never receives a forward or backward pass. We initially wrapped attention projections along with the MLP sublayers in Section 9.1’s ViT-B/16 experiment, and only caught the bypass by checking gradient flow into the attention-side LoRA parameters directly (all zero), rather than trusting that instantiating the wrapper was sufficient. Restricting LoRA to the plain `nn.Linear` layers inside each transformer block’s MLP sublayer avoids the issue entirely, since `nn.Linear.forward()` is called normally; this

is the configuration used to produce Table 6 (737,280 LoRA parameters out of 87,304,936 total, 0.84%).

D Diagnostic Data and Code

Raw per-run results, the sigmoid/bootstrap/hierarchical-model analysis scripts, and the full experiment campaign code (architecture zoo, task construction, training engine, Kaggle orchestration) are released with this submission at <https://github.com/kavyabhand/Self-Stabilizing-ML-Inference>.